

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn as sk
```

```
In [2]: df = pd.read_csv('credit_risk_dataset.csv')
```

```
In [3]: df
```

Out[3]:

	person_age	person_income	person_home_ownership	person_emp_length	
0	22	59000	RENT	123.0	
1	21	9600	OWN	5.0	
2	25	9600	MORTGAGE	1.0	
3	23	65500	RENT	4.0	
4	24	54400	RENT	8.0	
...	
32576	57	53000	MORTGAGE	1.0	
32577	54	120000	MORTGAGE	4.0	
32578	65	76000	RENT	3.0	HOMEIMP
32579	56	150000	MORTGAGE	5.0	
32580	66	42000	RENT	2.0	

32581 rows × 12 columns

```
In [4]: df.shape
```

Out[4]: (32581, 12)

```
In [7]: df.head()
```

Out[7]:

	person_age	person_income	person_home_ownership	person_emp_length	loan_intent	loan_status
0	22	59000	RENT	123.0	PERSONAL	PAID
1	21	9600	OWN	5.0	EDUCATION	PAID
2	25	9600	MORTGAGE	1.0	MEDICAL	PAID
3	23	65500	RENT	4.0	MEDICAL	PAID
4	24	54400	RENT	8.0	MEDICAL	PAID

In [8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32581 entries, 0 to 32580
Data columns (total 12 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   person_age                             32581 non-null  int64
 1   person_income                           32581 non-null  int64
 2   person_home_ownership                  32581 non-null  object
 3   person_emp_length                       31686 non-null  float64
 4   loan_intent                             32581 non-null  object
 5   loan_grade                             32581 non-null  object
 6   loan_amnt                              32581 non-null  int64
 7   loan_int_rate                           29465 non-null  float64
 8   loan_status                             32581 non-null  int64
 9   loan_percent_income                    32581 non-null  float64
10   cb_person_default_on_file              32581 non-null  object
11   cb_person_cred_hist_length             32581 non-null  int64
dtypes: float64(3), int64(5), object(4)
memory usage: 3.0+ MB
```

In [9]: `data_rows = df.shape[0]`
`data_columns = df.shape[1]`
`print(f'This dataset have {data_rows} rows and {data_columns} column`

This dataset have 32581 rows and 12 columns.

In [10]: `df.isnull().sum()`

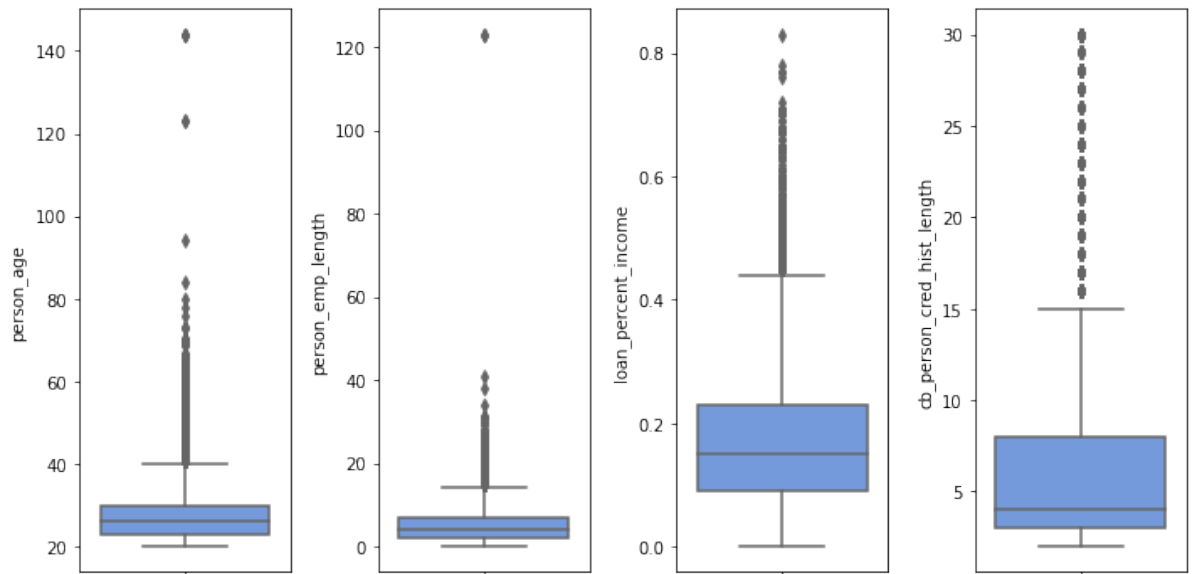
```
Out[10]: person_age                0
person_income                0
person_home_ownership        0
person_emp_length            895
loan_intent                   0
loan_grade                   0
loan_amnt                    0
loan_int_rate                 3116
loan_status                   0
loan_percent_income           0
cb_person_default_on_file     0
cb_person_cred_hist_length    0
dtype: int64
```

In [11]: `df.describe()`

Out[11]:

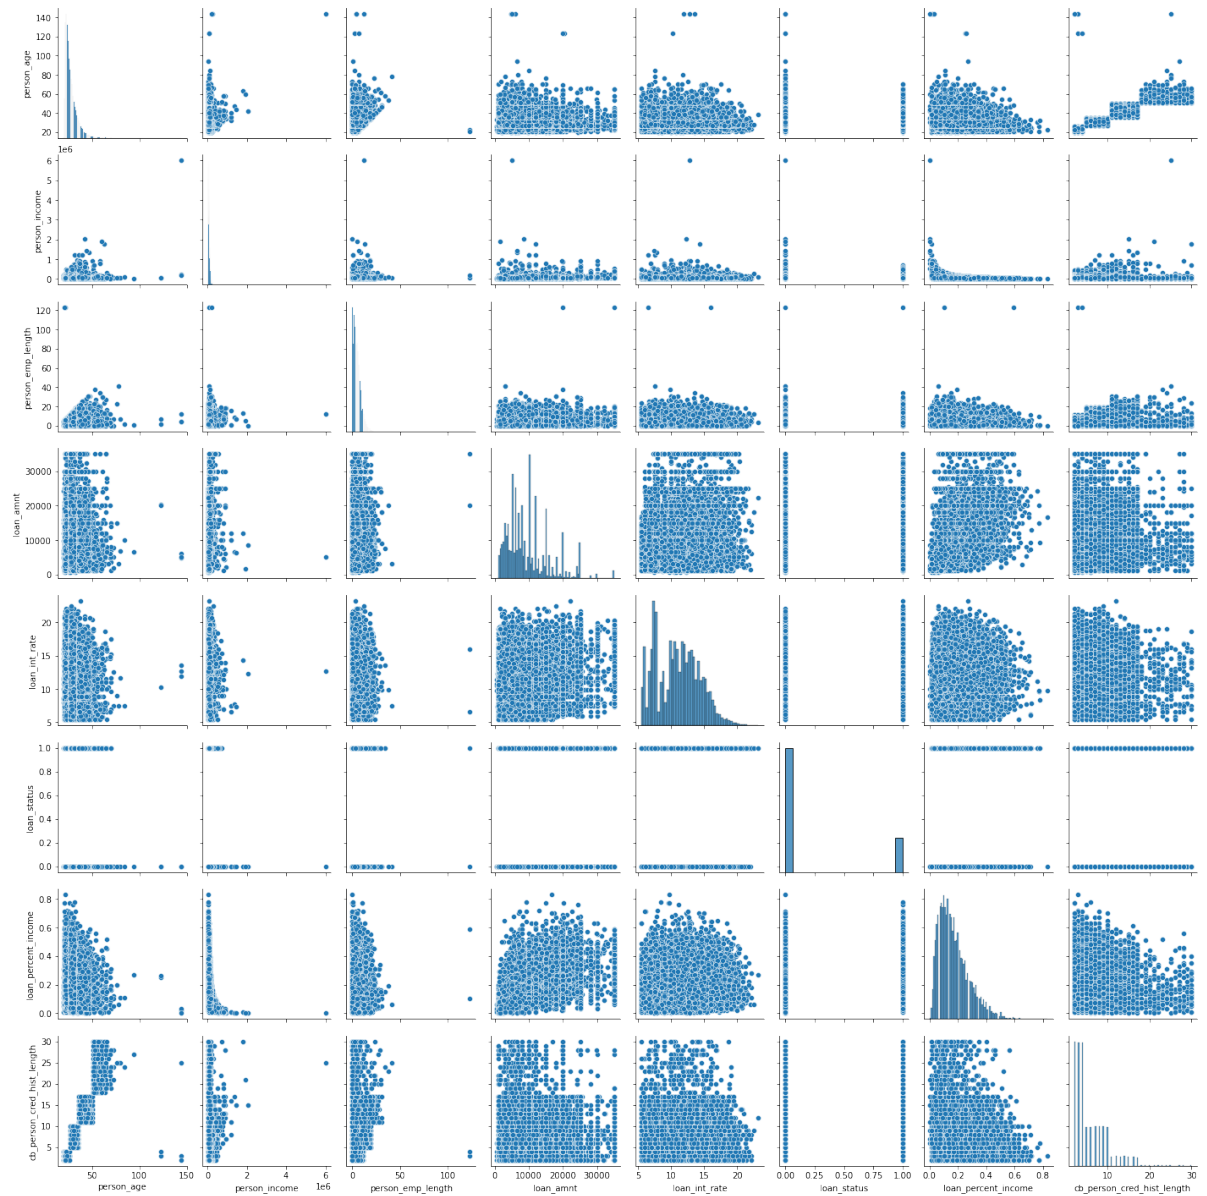
	person_age	person_income	person_emp_length	loan_amnt	loan_int_rate	loan_percent_income
count	32581.000000	3.258100e+04	31686.000000	32581.000000	29465.000000	32581.000000
mean	27.734600	6.607485e+04	4.789686	9589.371106	11.011695	15.012147
std	6.348078	6.198312e+04	4.142630	6322.086646	3.240459	12.124541
min	20.000000	4.000000e+03	0.000000	500.000000	5.420000	0.000000
25%	23.000000	3.850000e+04	2.000000	5000.000000	7.900000	5.000000
50%	26.000000	5.500000e+04	4.000000	8000.000000	10.990000	10.000000
75%	30.000000	7.920000e+04	7.000000	12200.000000	13.470000	15.000000
max	144.000000	6.000000e+06	123.000000	35000.000000	23.220000	100.000000

```
In [12]: features = ['person_age', 'person_emp_length', 'loan_percent_income',
plt.figure(figsize=(10,5))
for i in range(0,len(features)):
    plt.subplot(1, len(features), i + 1)
    sns.boxplot(y=df[features[i]], color='CornflowerBlue', orient='v')
    plt.tight_layout()
```



In [13]: `sns.pairplot(df)`

Out[13]: `<seaborn.axisgrid.PairGrid at 0x7fdc6b67b5e0>`




In [14]: `df.loan_status.value_counts()`

Out[14]:


0	25473
1	7108

Name: loan_status, dtype: int64

```
In [15]: plt.figure(figsize=(8,4))  
sns.countplot(x='loan_status', data=df)  
plt.show()
```



```
In [16]: plt.figure(figsize=(10,5))  
sns.histplot(data=df, x='person_age', bins=15)  
plt.show()
```




```
In [3]: variables = ['person_age', 'person_income', 'person_emp_length', 'loan_amnt', 'loan_status']  
credit_risk_corr = df[variables].corr()  
credit_risk_corr
```

Out[3]:

	person_age	person_income	person_emp_length	loan_amnt	loan_status
person_age	1.000000	0.173202	0.163106	0.050787	0.021629
person_income	0.173202	1.000000	0.134268	0.266820	-0.144449
person_emp_length	0.163106	0.134268	1.000000	0.113082	-0.082489
loan_amnt	0.050787	0.266820	0.113082	1.000000	0.105376
loan_int_rate	0.012580	0.000792	-0.056405	0.146813	0.041967
loan_percent_income	-0.042411	-0.254471	-0.054111	0.572612	0.859133
cb_person_cred_hist_length	0.859133	0.117987	0.144699	0.041967	1.000000
loan_status	-0.021629	-0.144449	-0.082489	0.105376	1.000000

```
In [4]: plt.figure(figsize=(8,8))
sns.heatmap(credit_risk_corr, cmap='Blues', annot=True, fmt='.2f')
plt.show()
```



```
In [5]: cat_cols = ['person_home_ownership', 'loan_intent', 'loan_grade', '
for i in cat_cols:
    print(f'Total row of variable {i}')
    print(df[i].value_counts())
    print()
```

Total row of variable person_home_ownership

RENT	16446
------	-------

MORTGAGE	13444
----------	-------

OWN	2584
-----	------

OTHER	107
-------	-----

Name: person_home_ownership, dtype: int64

Total row of variable loan_intent

EDUCATION	6453
-----------	------

MEDICAL	6071
---------	------

VENTURE	5719
---------	------

PERSONAL	5521
----------	------

DEBTCONSOLIDATION	5212
-------------------	------

HOMEIMPROVEMENT	3605
-----------------	------

Name: loan_intent, dtype: int64

Total row of variable loan_grade

A	10777
---	-------

B	10451
---	-------

C	6458
---	------

D	3626
---	------

E	964
---	-----

F	241
---	-----

G	64
---	----

Name: loan_grade, dtype: int64


Total row of variable cb_person_default_on_file

N	26836
---	-------

Y	5745
---	------

Name: cb_person_default_on_file, dtype: int64


```
In [6]: plt.figure(figsize=(15,9))
for i in range(0,len(cat_cols)):
    plt.subplot(2,2,i + 1)
    sns.countplot(data= df, x = cat_cols[i], hue='loan_status')
    plt.tight_layout()
```



```
In [7]: df.drop(df.loc[df['person_emp_length'] == 123].index, inplace=True)
```

```
In [8]: df.drop(df.loc[df['person_age'] >= 123].index, inplace=True)
```

```
In [9]: df.loc[df['person_age'] >= 123].index
```

```
Out[9]: Int64Index([], dtype='int64')
```

```
In [10]: num_cols = pd.DataFrame(df[df.select_dtypes(include=['float', 'int']
```

```
In [15]: num_cols_hist = num_cols.drop(['loan_status'], axis=1)
plt.figure(figsize=(12,16))

for i, col in enumerate(num_cols_hist.columns):
    idx = int('42'+ str(i+1))
    plt.subplot(idx)
    sns.distplot(num_cols_hist[col], color='forestgreen',
                  kde_kws={'color': 'indianred', 'lw': 2, 'label': ''})
    plt.title(col+' distribution', fontsize=14)
    plt.ylabel('Probablity', fontsize=12)
    plt.xlabel(col, fontsize=12)
    plt.xticks(fontsize=12)
    plt.yticks(fontsize=12)
    plt.legend(['KDE'], prop={"size":12})
```

```
plt.subplots_adjust(top=0.92, bottom=0.08, left=0.10, right=0.95, h
                    wspace=0.35)
plt.show()
```

/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



warnings.warn(msg, FutureWarning)

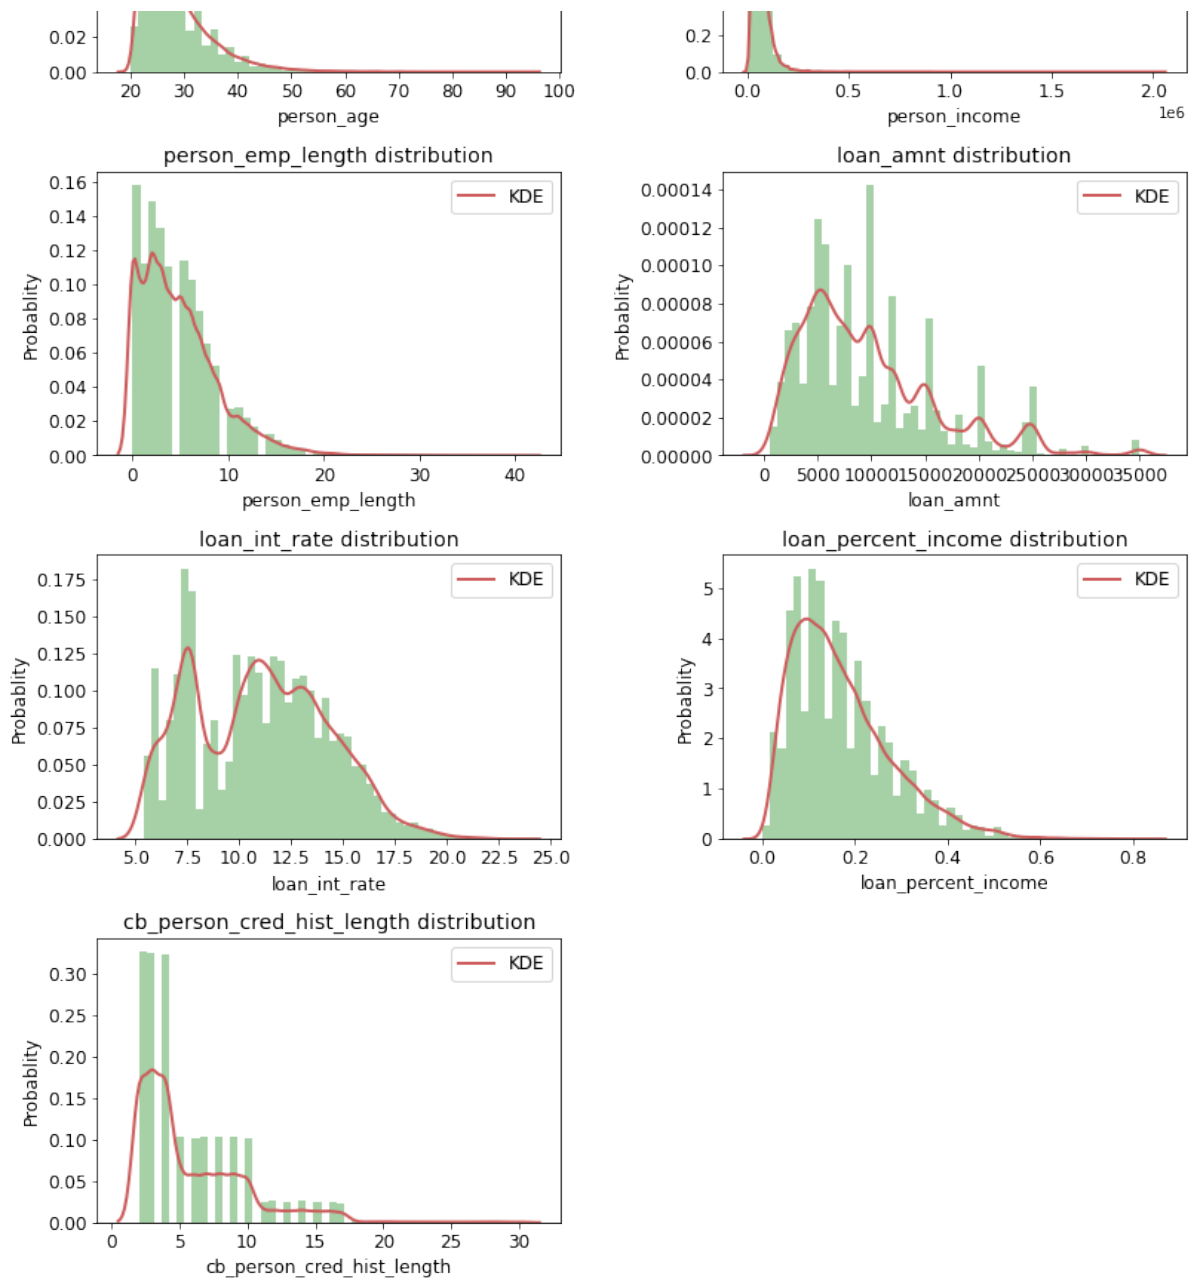
/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)





```
In [11]: df = df[df['person_age']<=100]
df = df[df['person_emp_length']<=60]
df = df[df['person_income']<=4e6]
```

```
In [12]: mean_person_emp_length = df['person_emp_length'].mean()
mean_loan_int_rate = df['loan_int_rate'].mean()

df['person_emp_length'] = df['person_emp_length'].fillna(mean_person_emp_length)
df['loan_int_rate'] = df['loan_int_rate'].fillna(mean_loan_int_rate)
```

```
In [13]: from sklearn.preprocessing import LabelEncoder
```

```
In [14]: label_encoder = LabelEncoder()
df["loan_intent"] = label_encoder.fit_transform(df["loan_intent"])
df["person_home_ownership"] = label_encoder.fit_transform(df["person_
df["loan_grade"] = label_encoder.fit_transform(df["loan_grade"])
df["cb_person_default_on_file"] = label_encoder.fit_transform(df["cb_
```

```
In [15]: df.head()
```

```
Out[15]:
```

	person_age	person_income	person_home_ownership	person_emp_length	loan_intent	lc
1	21	9600	2	5.0	1	
2	25	9600	0	1.0	3	
3	23	65500	3	4.0	3	
4	24	54400	3	8.0	3	
5	21	9900	2	2.0	5	

```
In [16]: df["person_age"].max()
```

```
Out[16]: 94
```

```
In [17]: df.isnull().sum()
```

```
Out[17]: person_age      0
person_income      0
person_home_ownership      0
person_emp_length      0
loan_intent      0
loan_grade      0
loan_amnt      0
loan_int_rate      0
loan_status      0
loan_percent_income      0
cb_person_default_on_file      0
cb_person_cred_hist_length      0
dtype: int64
```

```
In [18]: from scipy.stats import uniform, randint
from sklearn import model_selection, linear_model, metrics
from sklearn.metrics import auc, accuracy_score, confusion_matrix,
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from xgboost import XGBClassifier
from sklearn.metrics import plot_confusion_matrix
```

```
In [19]: R_tree=RandomForestClassifier()
lg = LogisticRegression(random_state=42)
N_bayes=GaussianNB()
XGB_model = XGBClassifier(learning_rate=0.1, max_depth=10, scale_po
knn = KNeighborsClassifier(n_neighbors=150)
```

```
In [20]: label = df['loan_status'] # labels
features = df.drop('loan_status',axis=1) # features
x_train, x_test, y_train, y_test = model_selection.train_test_split

print('The train dataset has {} data\nThe test dataset has {} data'
      format(x_train.shape[0], x_test.shape[0]))
```

The train dataset has 22175 data
The test dataset has 9504 data

```
In [24]: from sklearn.model_selection import cross_val_score
scores = cross_val_score(R_tree, x_train, y_train, cv = 10, scoring
print('Cross-validation scores:{}'.format(scores))
```

Cross-validation scores:[0.9220018 0.93688007 0.93462579 0.932371
51 0.93146979 0.9300857
0.932341 0.93279206 0.93053676 0.92061344]

```
In [25]: scores = cross_val_score(knn, x_train, y_train, cv = 10, scoring='a
print('Cross-validation scores:{}'.format(scores))
```

Cross-validation scores:[0.83453562 0.8367899 0.83769161 0.840396
75 0.83047791 0.83942264
0.83400992 0.8488949 0.83897158 0.8285972]

```
In [24]: scores = cross_val_score(N_bayes, x_train, y_train, cv = 10, scorin
print('Cross-validation scores:{}'.format(scores))
```

Cross-validation scores:[0.81514878 0.81875564 0.82506763 0.816501
35 0.82100992 0.82318448
0.82995038 0.81145692 0.81686964 0.81867388]

```
In [26]: scores = cross_val_score(lg, x_train, y_train, cv = 10, scoring='ac
print('Cross-validation scores:{}'.format(scores))
```

Cross-validation scores:[0.80342651 0.8097385 0.80297565 0.812443
64 0.80613165 0.81235904
0.81326116 0.8105548 0.8015336 0.80514208]

```
In [29]: scores = cross_val_score(XGB_model, x_train, y_train, cv = 10, scor
print('Cross-validation scores:{}'.format(scores))
```

```
/opt/anaconda3/lib/python3.8/site-packages/xgboost/sklearn.py:1224
: UserWarning: The use of label encoder in XGBClassifier is deprec
ated and will be removed in a future release. To remove this warni
ng, do the following: 1) Pass option use_label_encoder=False when
constructing XGBClassifier object; and 2) Encode your labels (y) a
s integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
Cross-validation scores:[0.93327322 0.9418395 0.93958521 0.935527
5 0.94274121 0.9391069
0.93594948 0.94091114 0.92918358 0.92737934]
```

```
In [21]: def model_assess(model, name='Default'):
model.fit(x_train, y_train)
preds = model.predict(x_test)
preds_proba = model.predict_proba(x_test)
print(name, '\n',classification_report(y_test, model.predict(x_
print(confusion_matrix(y_test,preds))
```

```
In [22]: N_bayes=GaussianNB()
model_assess(N_bayes, name='Naive bayes')

lg = LogisticRegression(random_state=42)
model_assess(lg, 'Logistic Regression')

R_tree=RandomForestClassifier()
model_assess(R_tree,'RandomForest Classifier')

XGB_model = XGBClassifier(learning_rate=0.1, max_depth=10, scale_po
model_assess(XGB_model, 'Xgboost')

knn = KNeighborsClassifier(n_neighbors=150)
model_assess(knn, name='KNN')
```

Naive bayes

	precision	recall	f1-score	support
0	0.85	0.96	0.90	7456
1	0.70	0.37	0.48	2048
accuracy			0.83	9504
macro avg	0.77	0.66	0.69	9504
weighted avg	0.81	0.83	0.81	9504

```
[[7135 321]
 [1296 752]]
Logistic Regression
```

	precision	recall	f1-score	support
0	0.81	0.98	0.89	7456
1	0.71	0.15	0.25	2048
accuracy			0.80	9504
macro avg	0.76	0.57	0.57	9504
weighted avg	0.79	0.80	0.75	9504

```
[[7327 129]
 [1738 310]]
RandomForest Classifier
```

	precision	recall	f1-score	support
0	0.93	0.99	0.96	7456
1	0.97	0.71	0.82	2048
accuracy			0.93	9504
macro avg	0.95	0.85	0.89	9504
weighted avg	0.94	0.93	0.93	9504

```
[[7411 45]
 [ 584 1464]]
```

/opt/anaconda3/lib/python3.8/site-packages/xgboost/sklearn.py:1224
: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
warnings.warn(label_encoder_deprecation_msg, UserWarning)

```
Xgboost
```

	precision	recall	f1-score	support
0	0.93	0.99	0.96	7456
1	0.95	0.75	0.84	2048
accuracy			0.94	9504
macro avg	0.94	0.87	0.90	9504
weighted avg	0.94	0.94	0.93	9504

```
[[7373 83]
 [ 517 1531]]
KNN
```

	precision	recall	f1-score	support
0	0.85	0.96	0.90	7456
1	0.71	0.38	0.49	2048
accuracy			0.83	9504

```

macro avg      0.78      0.67      0.70      9504
weighted avg   0.82      0.83      0.81      9504

[[7144  312]
 [1276  772]]

```

```
In [23]: from tabulate import tabulate
```

```
In [25]: head = ["Model", "Accuracy", "Precision", "Recall", "F1-score"]

final_metrics = [
    ["Naive Bayes", 0.83, 0.81, 0.83, 0.81],
    ["Logistic Regression", 0.80, 0.79, 0.80, 0.75],
    ["KNN Classifier", 0.83, 0.82, 0.83, 0.81],
    ["RandomForestClassifier", 0.93, 0.94, 0.93, 0.93],
    ["Xgboost", 0.94, 0.94, 0.94, 0.93]
]

print(tabulate(final_metrics, headers = head))
```

Model	Accuracy	Precision	Recall	F1-score
Naive Bayes	0.83	0.81	0.83	0.81
Logistic Regression	0.8	0.79	0.8	0.75
KNN Classifier	0.83	0.82	0.83	0.81
RandomForestClassifier	0.93	0.94	0.93	0.93
Xgboost	0.94	0.94	0.94	0.93

```
In [31]: ### ROC AUC
fig = plt.figure(figsize=(8,5))
plt.plot([0, 1], [0, 1], 'r--')

preds_proba_knn = knn.predict_proba(x_test)
probsknn = preds_proba_knn[:, 1]
fpr, tpr, thresh = metrics.roc_curve(y_test, probsknn)
aucknn = roc_auc_score(y_test, probsknn)
plt.plot(fpr, tpr, label=f'KNN, AUC = {str(round(aucknn,3))}')

preds_proba_N_bayes = N_bayes.predict_proba(x_test)
probsknn = preds_proba_knn[:, 1]
fpr, tpr, thresh = metrics.roc_curve(y_test, probsknn)
aucknn = roc_auc_score(y_test, probsknn)
plt.plot(fpr, tpr, label=f'Naive Bayes, AUC = {str(round(aucknn,3))}')

#Logistic Regression
```



```


#Logistic Regression
preds_proba_lg = lg.predict_proba(x_test)
probslg = preds_proba_lg[:, 1]
fpr, tpr, thresh = metrics.roc_curve(y_test, probslg)
auc_lg = roc_auc_score(y_test, probslg)
plt.plot(fpr, tpr, label=f'Logistic Regression, AUC = {str(round(auc_lg,3))}')

#RandomForest Classifier
preds_proba_R_tree = R_tree.predict_proba(x_test)
probsR_tree = preds_proba_R_tree[:, 1]
fpr, tpr, thresh = metrics.roc_curve(y_test, probsR_tree)
auc_lg = roc_auc_score(y_test, probsR_tree)
plt.plot(fpr, tpr, label=f'RandomForest Classifier, AUC = {str(round(auc_lg,3))}')

#XGBoost
preds_proba_xgb = XGB_model.predict_proba(x_test)
probsxgb = preds_proba_xgb[:, 1]
fpr, tpr, thresh = metrics.roc_curve(y_test, probsxgb)
aucxgb = roc_auc_score(y_test, probsxgb)
plt.plot(fpr, tpr, label=f'XGBoost, AUC = {str(round(aucxgb,3))}')

plt.ylabel("True Positive Rate", fontsize=12)
plt.xlabel("False Positive Rate", fontsize=12)
plt.title("ROC curve")
plt.rcParams['axes.titlesize'] = 16
plt.legend()
plt.show()

```




In []:

```
In [44]: from sklearn.metrics import confusion_matrix  
from sklearn.metrics import plot_confusion_matrix
```


```
In [55]: plot_confusion_matrix(N_bayes, x_test, y_test)
```

```
Out [55]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at  
0x7f89205950a0>
```




```
In [47]: plot_confusion_matrix(knn, x_test, y_test)
```

```
Out [47]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at  
0x7f88f016c190>
```




```
In [53]: plot_confusion_matrix(lg, x_test, y_test)
```

```
Out [53]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f8930dd43d0>
```




```
In [54]: plot_confusion_matrix(R_tree, x_test, y_test)
```

```
Out [54]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f88e039a9a0>
```



```
In [52]: plot_confusion_matrix(XGB_model, x_test, y_test)
```

```
Out[52]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f8931f3af40>
```



```
In [ ]:
```

```
In [ ]: # KNeighbors Classifier
```

```
In [20]: grid_params = { 'n_neighbors' : [5,7,9,11,13,15],  
                        'weights' : ['uniform','distance'],  
                        'metric' : ['minkowski','euclidean','manhattan']}
```

```
In [22]: from sklearn.model_selection import GridSearchCV  
gs = GridSearchCV(KNeighborsClassifier(), grid_params, verbose = 1,
```

```
In [24]: gs.fit(x_train,y_train)
print(f"Best Score: {gs.best_score_}")
print("Standard Devaition:",gs.cv_results_['std_test_score'][gs.bes
print("Best parameters set:")
best_parameters = gs.best_estimator_.get_params()
for param_name in sorted(grid_params.keys()):
    print(f"\t{param_name}: {best_parameters[param_name]}")
```

Fitting 3 folds for each of 36 candidates, totalling 108 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 2.5s

Best Score: 0.8443741381478156

Standard Devaition: 0.0026381041572729655

Best parameters set:

metric: manhattan

n_neighbors: 15

weights: distance

[Parallel(n_jobs=-1)]: Done 108 out of 108 | elapsed: 4.1s finished

```
In [25]: g_res = gs.fit(x_train, y_train)
```

Fitting 3 folds for each of 36 candidates, totalling 108 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 52 tasks | elapsed: 1.4s

[Parallel(n_jobs=-1)]: Done 93 out of 108 | elapsed: 2.4s remaining: 0.4s

[Parallel(n_jobs=-1)]: Done 108 out of 108 | elapsed: 2.6s finished

```
In [26]: g_res.best_score_
```

```
Out[26]: 0.8443741381478156
```

```
In [27]: g_res.best_params_
```

```
Out[27]: {'metric': 'manhattan', 'n_neighbors': 15, 'weights': 'distance'}
```

```
In [29]: gs.cv_results_
5320168, 0.2270871 ,
        0.0407571 , 0.24188296, 0.05830042, 0.25471973, 0.06716299
,
        0.24916704, 0.07165273, 0.24579279, 0.0480605 , 0.23538852
,
        0.06678303, 0.302713 , 0.07904553, 0.32016063, 0.06641014
,
        0.28328514, 0.06828213, 0.26133839, 0.0660344 , 0.21840445
,
        0.05286956, 0.23570021, 0.05055467, 0.23927943, 0.05580799
,
        0.22160697, 0.07106972, 0.24949837, 0.06862577, 0.18994912
,
        0.06747842]),
'std_score_time': array([0.00822901, 0.00514166, 0.01411297, 0.00
735241, 0.02825148,
        0.00309762, 0.03853331, 0.00211289, 0.03971449, 0.00536478
,
        0.01180724, 0.01513796, 0.02797528, 0.00166903, 0.04598741
```

```
In [30]: result=pd.DataFrame(gs.cv_results_)
print(result)

   mean_fit_time  std_fit_time  mean_score_time  std_score_time p
arametric \
0      0.019805      0.000917      0.240762      0.008229
minkowski
1      0.022732      0.002807      0.046491      0.005142
minkowski
2      0.018147      0.005314      0.225082      0.014113
minkowski
3      0.015069      0.000466      0.053202      0.007352
minkowski
4      0.014763      0.002351      0.227087      0.028251
minkowski
5      0.013540      0.002371      0.040757      0.003098
minkowski
6      0.017636      0.003851      0.241883      0.038533
minkowski
7      0.017406      0.003491      0.058300      0.002113
minkowski
8      0.015601      0.000580      0.254720      0.039714
minkowski
```

In []:

In []:

In []:

```
In [ ]: # Logistic Regression
```

```
In [31]: LR_param_grid = {
          'C': np.logspace(0, 4, num=10),
          'penalty': ['l2'],
          'solver': ['liblinear', 'saga', 'newton-cg', 'lbfgs']
        }
```

```
In [32]: logistic = LogisticRegression()

Logistic_model = model_selection.RandomizedSearchCV(
    estimator = logistic,
    param_distributions = LR_param_grid,
    n_iter = 20,
    scoring = "accuracy",
    verbose = 5,
    n_jobs = 1,
    cv = 5
)
```

```
In [33]: Logistic_model.fit(x_train,y_train)
print(f"Best Score: {Logistic_model.best_score}")
print("Standard Deviation:",Logistic_model.cv_results_['std_test_sc
print("Best parameters set:")
best_parameters = Logistic_model.best_estimator_.get_params()
for param_name in sorted(LR_param_grid.keys()):
print(f"\t{param_name}: {best_parameters[param_name]}")
```

```
Fitting 5 folds for each of 20 candidates, totalling 100 fits
[CV] solver=newton-cg, penalty=l2, C=21.544346900318832 .....
....
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concu
rrent workers.
/opt/anaconda3/lib/python3.8/site-packages/scipy/optimize/linesear
ch.py:477: LineSearchWarning: The line search algorithm did not co
nverge
  warn('The line search algorithm did not converge', LineSearchWar
ning)
/opt/anaconda3/lib/python3.8/site-packages/scipy/optimize/linesear
ch.py:327: LineSearchWarning: The line search algorithm did not co
nverge
  warn('The line search algorithm did not converge', LineSearchWar
ning)
/opt/anaconda3/lib/python3.8/site-packages/sklearn/utils/optimize.
py:211: ConvergenceWarning: newton-cg failed to converge. Increase
the number of iterations.
  warnings.warn("newton-cg failed to converge. Increase the "
```

```
In [34]: Logistic_model.cv_results_
```

```
Out[34]: {'mean_fit_time': array([0.60576491, 0.01789865, 0.20223923, 0.200
```

```

76528, 0.64459615,
      0.01796274, 0.63561707, 0.20858283, 0.20496206, 0.04289246
',
      0.04218559, 0.60432382, 0.65315065, 0.04150758, 0.01874418
',
      0.04258528, 0.01881323, 0.02018046, 0.20426989, 0.20889564
]),
'std_fit_time': array([0.04925154, 0.00227371, 0.0021942 , 0.0020
789 , 0.01315948,
      0.00142151, 0.03296847, 0.00603425, 0.00143005, 0.00682262
',
      0.00641753, 0.00899923, 0.0311405 , 0.00569562, 0.00133522
',
      0.00694027, 0.00182796, 0.00201667, 0.00283887, 0.00290232
]),
'mean_score_time': array([0.00159802, 0.00116835, 0.00187278, 0.0
0176234, 0.00157585,
      0.00117164, 0.00139194, 0.00192318, 0.00171061, 0.00137129
',
      0.00137019, 0.00149293, 0.00152178, 0.00137038, 0.00121737
',
      0.00137825, 0.00117159, 0.00139918, 0.00151815, 0.00189476
]),
'std_score_time': array([1.09841134e-04, 2.02622691e-05, 2.515765
46e-04, 2.29276446e-04,
      2.49128802e-04, 9.44209096e-05, 7.00354160e-05, 2.87488978
e-04,
      2.51206414e-04, 8.13724532e-05, 8.94229239e-05, 1.02717343
e-04,
      2.56664545e-04, 6.84694658e-05, 2.94949897e-05, 1.35802218
e-04,
      7.68996573e-05, 2.30319534e-04, 1.24908911e-04, 2.25424604
e-04]),
'param_solver': masked_array(data=['newton-cg', 'lbfgs', 'saga',
'saga', 'newton-cg',
      'lbfgs', 'newton-cg', 'saga', 'saga', 'libline
ar',
      'liblinear', 'newton-cg', 'newton-cg', 'liblin
ear',
      'lbfgs', 'liblinear', 'lbfgs', 'lbfgs', 'saga'
, 'saga'],
      mask=[False, False, False, False, False, False, Fals
e, False,
      False, False, False, False, False, False, Fals
e, False,
      False, False, False, False],
      fill_value='?',
      dtype=object),
'param_penalty': masked_array(data=['l2', 'l2', 'l2', 'l2', 'l2',
'l2', 'l2', 'l2', 'l2',
      'l2', 'l2', 'l2', 'l2', 'l2', 'l2', 'l2', 'l2'
, 'l2',
      'l2', 'l2'],

```



```

        mask=[False, False, False, False, False, False, False, False,
e, False,
        False, False, False, False, False, False, False, False,
e, False,
        False, False, False, False],
        fill_value='?',
        dtype=object),
'param_C': masked_array(data=[21.544346900318832, 3593.8136638046
26,
        21.544346900318832, 1291.5496650148827,
        1291.5496650148827, 21.544346900318832,
        464.15888336127773, 1.0, 166.81005372000593,
        464.15888336127773, 1291.5496650148827,
        59.94842503189409, 3593.813663804626,
        21.544346900318832, 166.81005372000593,
        2.7825594022071245, 464.15888336127773, 1.0,
        7.742636826811269, 2.7825594022071245],
        mask=[False, False, False, False, False, False, False, False,
e, False,
        False, False, False, False, False, False, False, False,
e, False,
        False, False, False, False],
        fill_value='?',
        dtype=object),
'params': [{'solver': 'newton-cg', 'penalty': 'l2', 'C': 21.54434
6900318832},
        {'solver': 'lbfgs', 'penalty': 'l2', 'C': 3593.813663804626},
        {'solver': 'saga', 'penalty': 'l2', 'C': 21.544346900318832},
        {'solver': 'saga', 'penalty': 'l2', 'C': 1291.5496650148827},
        {'solver': 'newton-cg', 'penalty': 'l2', 'C': 1291.5496650148827
},
        {'solver': 'lbfgs', 'penalty': 'l2', 'C': 21.544346900318832},
        {'solver': 'newton-cg', 'penalty': 'l2', 'C': 464.15888336127773
},
        {'solver': 'saga', 'penalty': 'l2', 'C': 1.0},
        {'solver': 'saga', 'penalty': 'l2', 'C': 166.81005372000593},
        {'solver': 'liblinear', 'penalty': 'l2', 'C': 464.15888336127773
},
        {'solver': 'liblinear', 'penalty': 'l2', 'C': 1291.5496650148827
},
        {'solver': 'newton-cg', 'penalty': 'l2', 'C': 59.94842503189409}
,
        {'solver': 'newton-cg', 'penalty': 'l2', 'C': 3593.813663804626}
,
        {'solver': 'liblinear', 'penalty': 'l2', 'C': 21.544346900318832
},
        {'solver': 'lbfgs', 'penalty': 'l2', 'C': 166.81005372000593},
        {'solver': 'liblinear', 'penalty': 'l2', 'C': 2.7825594022071245
},
        {'solver': 'lbfgs', 'penalty': 'l2', 'C': 464.15888336127773},
        {'solver': 'lbfgs', 'penalty': 'l2', 'C': 1.0},
        {'solver': 'saga', 'penalty': 'l2', 'C': 7.742636826811269},
        {'solver': 'saga', 'penalty': 'l2', 'C': 2.7825594022071245}],

```

```

'split0_test_score': array([0.8410372 , 0.8065389 , 0.80631342, 0
.80631342, 0.84126268,
      0.8065389 , 0.84193912, 0.80631342, 0.80631342, 0.82322435
,
      0.82390079, 0.84239008, 0.84193912, 0.82322435, 0.8065389
,
      0.82322435, 0.8065389 , 0.8065389 , 0.80631342, 0.80631342
]),
'split1_test_score': array([0.8475761 , 0.80766629, 0.80586246, 0
.80586246, 0.84735062,
      0.80766629, 0.84712514, 0.80586246, 0.80586246, 0.82998873
,
      0.82998873, 0.84870349, 0.84712514, 0.82998873, 0.80766629
,
      0.82998873, 0.80766629, 0.80766629, 0.80586246, 0.80586246
]),
'split2_test_score': array([0.84870349, 0.80947012, 0.81014656, 0
.81014656, 0.84847802,
      0.80947012, 0.84847802, 0.81014656, 0.81014656, 0.82908681
,
      0.81894025, 0.84847802, 0.84870349, 0.81894025, 0.80947012
,
      0.82931229, 0.80947012, 0.80947012, 0.81014656, 0.81014656
]),
'split3_test_score': array([0.8518602 , 0.81172492, 0.81059752, 0
.81059752, 0.8518602 ,
      0.81172492, 0.8518602 , 0.81059752, 0.81059752, 0.80473506
,
      0.80473506, 0.85163472, 0.8518602 , 0.80473506, 0.81172492
,
      0.80473506, 0.81172492, 0.81172492, 0.81059752, 0.81059752
]),
'split4_test_score': array([0.83878241, 0.80360767, 0.80180383, 0
.80180383, 0.83720406,
      0.80360767, 0.8378805 , 0.80180383, 0.80180383, 0.81894025
,
      0.81894025, 0.83742954, 0.8378805 , 0.81894025, 0.80360767
,
      0.81894025, 0.80360767, 0.80360767, 0.80180383, 0.80180383
]),
'mean_test_score': array([0.84559188, 0.80780158, 0.80694476, 0.8
0694476, 0.84523112,
      0.80780158, 0.8454566 , 0.80694476, 0.80694476, 0.82119504
,
      0.81930101, 0.84572717, 0.84550169, 0.81916573, 0.80780158
,
      0.82124014, 0.80780158, 0.80780158, 0.80694476, 0.80694476
]),
'std_test_score': array([0.00489949, 0.00273416, 0.00321227, 0.00
321227, 0.00527563,
      0.00273416, 0.00495357, 0.00321227, 0.00321227, 0.00916234
,
      0.00834041, 0.00512468, 0.00498181, 0.00826989, 0.00273416

```

```
,
      0.00920154, 0.00273416, 0.00273416, 0.00321227, 0.00321227
]),
'rank_test_score': array([ 2, 10, 15, 15,  5, 10,  4, 15, 15,  7,
 8,  1,  3,  9, 10,  6, 10,
10, 15, 15], dtype=int32)}
```

```
In [35]: result=pd.DataFrame(Logistic_model.cv_results_)
print(result)
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	p
aram_solver \					
0	0.605765	0.049252	0.001598	0.000110	
newton-cg					
1	0.017899	0.002274	0.001168	0.000020	
lbfgs					
2	0.202239	0.002194	0.001873	0.000252	
saga					
3	0.200765	0.002079	0.001762	0.000229	
saga					
4	0.644596	0.013159	0.001576	0.000249	
newton-cg					
5	0.017963	0.001422	0.001172	0.000094	
lbfgs					
6	0.635617	0.032968	0.001392	0.000070	
newton-cg					
7	0.208583	0.006034	0.001923	0.000287	
saga					
8	0.204962	0.001430	0.001711	0.000251	

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]: # Naive Bayes
```

```
In [33]: from sklearn.model_selection import RandomizedSearchCV
classifier2 = RandomizedSearchCV(R_tree,params,cv=5)
```

```
In [31]: params = {
'var_smoothing': np.logspace(0,-9, num=100)
}
```

```
In [34]: classifier4 = RandomizedSearchCV(N_bayes,params,cv=5)
```

In [35]: `classifier4.fit(x_train,y_train)`

Out [35]: RandomizedSearchCV(cv=5, estimator=GaussianNB(),
 param_distributions={'var_smoothing': array([1.
 00000000e+00, 8.11130831e-01, 6.57933225e-01, 5.33669923e-01,
 4.32876128e-01, 3.51119173e-01, 2.84803587e-01, 2.31012970e
 -01,
 1.87381742e-01, 1.51991108e-01, 1.23284674e-01, 1.00000000e
 -01,
 8.11130831e-02, 6.57933225e-02, 5.33669923e-02, 4.32876128e
 -02,
 3.51119173e-02, 2.84...
 1.23284674e-07, 1.00000000e-07, 8.11130831e-08, 6.57933225e
 -08,
 5.33669923e-08, 4.32876128e-08, 3.51119173e-08, 2.84803587e
 -08,
 2.31012970e-08, 1.87381742e-08, 1.51991108e-08, 1.23284674e
 -08,
 1.00000000e-08, 8.11130831e-09, 6.57933225e-09, 5.33669923e
 -09,
 4.32876128e-09, 3.51119173e-09, 2.84803587e-09, 2.31012970e
 -09,
 1.87381742e-09, 1.51991108e-09, 1.23284674e-09, 1.00000000e
 -09])))

In [36]: `classifier4.cv_results_`

Out [36]: {'mean_fit_time': array([0.00831461, 0.00435195, 0.00388036, 0.003
 46899, 0.0034637 ,
 0.00345373, 0.00345793, 0.00346546, 0.00343475, 0.00344124
]),
 'std_fit_time': array([3.59709244e-03, 3.73388572e-04, 7.00713065
 e-04, 1.57052527e-05,
 2.10842920e-05, 1.00487138e-05, 1.91464222e-05, 3.87259299
 e-05,
 9.38121819e-06, 1.49772418e-05]),
 'mean_score_time': array([0.00240235, 0.00149393, 0.00117865, 0.0
 0113735, 0.00112734,
 0.00112901, 0.00114446, 0.00113096, 0.00111995, 0.00112214
]),
 'std_score_time': array([4.61878940e-04, 2.06275147e-04, 7.022349
 60e-05, 7.44629737e-06,
 5.05222414e-06, 4.86186870e-06, 2.53878958e-05, 1.16553282
 e-05,
 5.67135611e-06, 4.37704305e-06]),
 'param_var_smoothing': masked_array(data=[5.3366992312063123e-05,
 0.12328467394420659,
 0.002848035868435802, 3.5111917342151273e-09,
 6.579332246575682e-07, 2.848035868435799e-06,
 8.111308307896872e-07, 1e-05, 0.00035111917342
 15131,
 1.519911082952933e-06],
 mask=[False, False, False, False, False, False, Fals

```

e, False,
                False, False],
                fill_value='?',
                dtype=object),
'params': [{'var_smoothing': 5.3366992312063123e-05},
{'var_smoothing': 0.12328467394420659},
{'var_smoothing': 0.002848035868435802},
{'var_smoothing': 3.5111917342151273e-09},
{'var_smoothing': 6.579332246575682e-07},
{'var_smoothing': 2.848035868435799e-06},
{'var_smoothing': 8.111308307896872e-07},
{'var_smoothing': 1e-05},
{'var_smoothing': 0.0003511191734215131},
{'var_smoothing': 1.519911082952933e-06}],
'split0_test_score': array([0.78624577, 0.78444194, 0.78692221, 0
.80202931, 0.78624577,
                0.78624577, 0.78624577, 0.78624577, 0.78737317, 0.78624577
]),
'split1_test_score': array([0.78759865, 0.78444194, 0.78782413, 0
.80856821, 0.78782413,
                0.78759865, 0.78782413, 0.78759865, 0.78759865, 0.78782413
]),
'split2_test_score': array([0.78872604, 0.78466742, 0.78692221, 0
.80225479, 0.789177 ,
                0.789177 , 0.789177 , 0.789177 , 0.78850056, 0.789177
]),
'split3_test_score': array([0.78308906, 0.78466742, 0.78827508, 0
.80225479, 0.78354002,
                0.78286359, 0.78354002, 0.78286359, 0.78466742, 0.78308906
]),
'split4_test_score': array([0.78714769, 0.78466742, 0.78624577, 0
.80315671, 0.78737317,
                0.78737317, 0.78737317, 0.78737317, 0.78759865, 0.78737317
]),
'mean_test_score': array([0.78656144, 0.78457723, 0.78723788, 0.8
0365276, 0.78683202,
                0.78665163, 0.78683202, 0.78665163, 0.78714769, 0.78674183
]),
'std_test_score': array([0.00191113, 0.00011046, 0.00072153, 0.00
248805, 0.0018951 ,
                0.0021123 , 0.0018951 , 0.0021123 , 0.0012992 , 0.00205372
]),
'rank_test_score': array([ 9, 10,  2,  1,  4,  7,  4,  7,  3,  6]
, dtype=int32)}

```

```

In [37]: best_parameters=classifier4.best_params_
print(best_parameters)

```

```

{'var_smoothing': 3.5111917342151273e-09}

```

```
In [38]: highest_accuracy=classifier4.best_score_  
print(highest_accuracy)
```

0.803652762119504

In []:

In []:

In []:

In []:

In []: